

1)

- A) T2 → T1  
It is conflict-serializable

T1	T2
	Write C
	Write A
Write D	
Read C	
Write A	

- B) It can
- |          |          |
|----------|----------|
| T1       | T2       |
| X-Lock D |          |
| Write D  |          |
|          | X-Lock C |
|          | Write C  |
|          | Unlock C |
| S-Lock C |          |
| Read C   |          |
|          | X-Lock A |
|          | Write A  |
|          | Unlock A |
|          | Commit   |
| X-Lock A |          |
| Write A  |          |
| Unlock D |          |
| Unlock A |          |
| Commit   |          |

- C) i. Yes. Read C in T1 would be done before Write C in T2 is committed  
 ii. It should only commit after T2 committed  
 iii. If T2 commit, T1 can also commit. If T2 abort, T1 should undo Write D
- D) If  $TS(T2) < TS(T1)$ , then the transaction would flow normally as T2 executes write C earlier than T1 executes read C and T2 executes write A before T1 executes write A

2)

- 1) False. The function of checkpoint is to log the modified pages (wrote) and to manage their rollbacks.
- 2) True. If the system crashed in the middle of a transaction, checkpoints provide logs allowing the administrator to decide actions on the transaction.
- 3) True. However its usefulness is dependent on the time between checkpoints
- 5) Then the manager have to continue from the last checkpoint with no previous damaged checkpoint.

3)

- a) Two of the log pages will be read. It will stop after reading T1 committed
- b) T0, T1, T4 will be redone. T2, T3, T5 will be undone.
- c) At the end of recovery,  
A = 100  
B = 800  
C = 2100